

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("10_USA_Housing.csv")
df
```

Out[2]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991- 3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

5000 rows × 7 columns

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

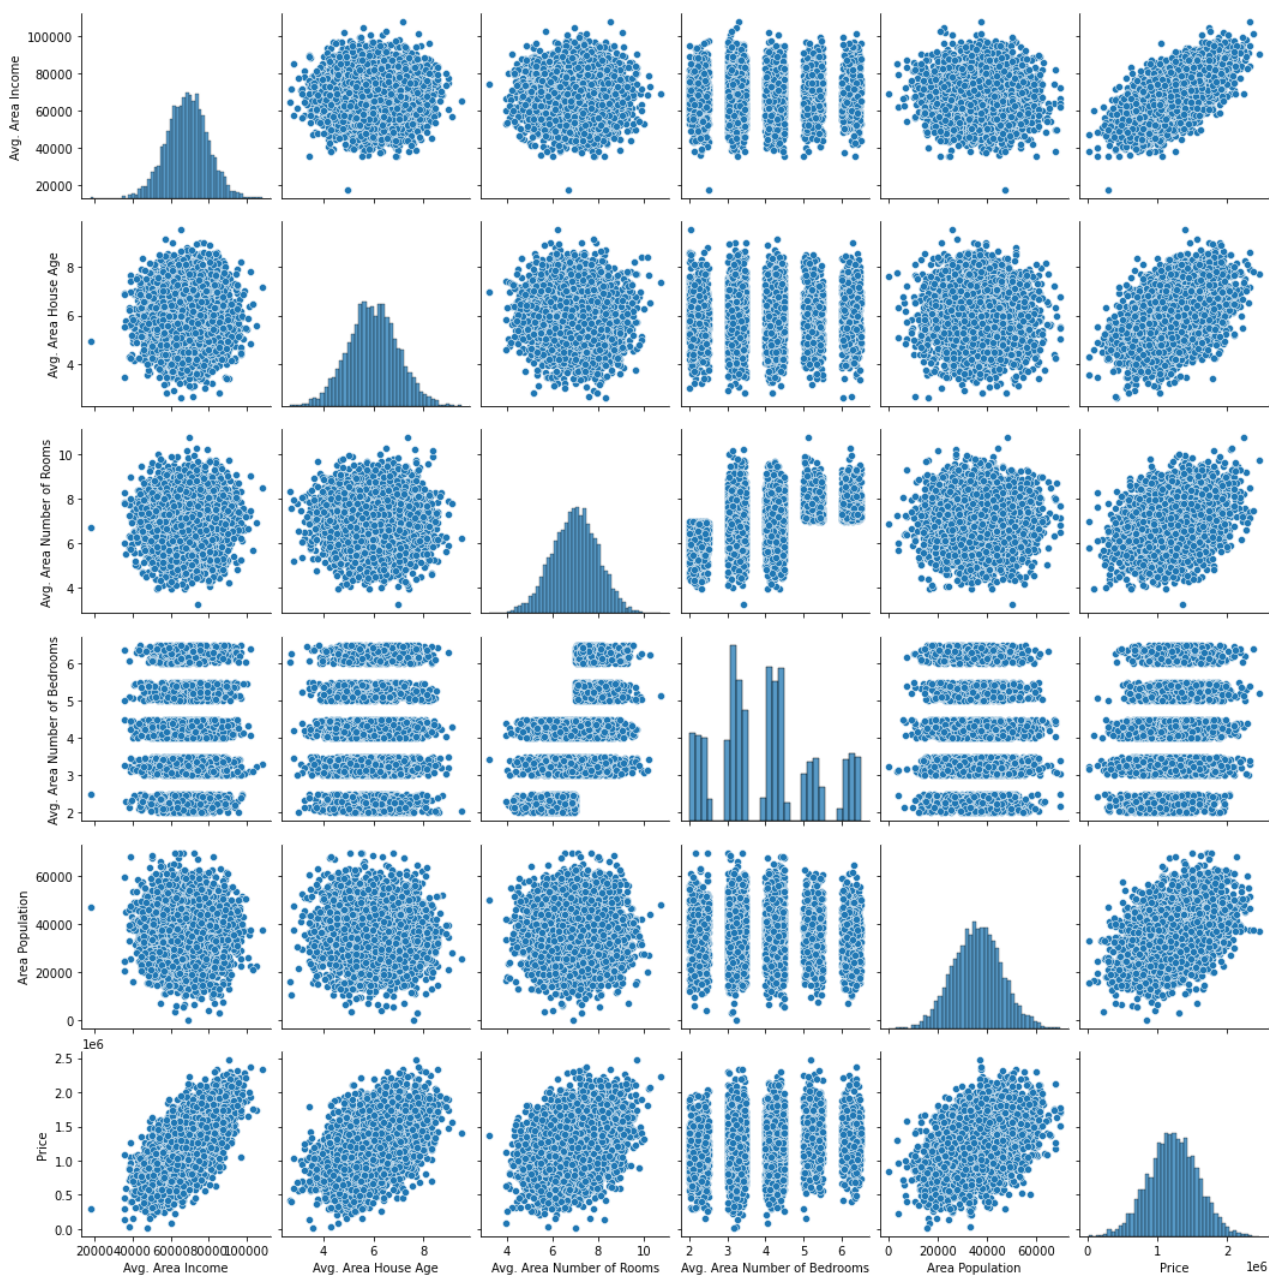
In [4]: df.describe()

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06


```
In [5]: sns.pairplot(df)
```

```
Out[5]: <seaborn.axisgrid.PairGrid at 0x1f2d5d1dd60>
```



```
In [6]: sns.displot(df['Price'])
```

```
Out[6]: <seaborn.axisgrid.FacetGrid at 0x1f2dbf8b760>
```



```
In [7]: df1=df.drop(['Address'],axis=1)
df1
```

```
Out[7]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06

5000 rows × 6 columns

```
In [8]: sns.heatmap(df1.corr())
```

```
Out[8]: <AxesSubplot:>
```



```
In [9]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression
```

```
In [10]: y=df['Price']  
x=df1.drop(['Price'],axis=1)  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [11]: model=LinearRegression()  
model.fit(x_train,y_train)  
model.intercept_
```

```
Out[11]: -2657812.6571027627
```


```
In [12]: coeff=pd.DataFrame(model.coef_,x.columns,columns=["Coefficient"])  
coeff
```

Out[12]:

	Coefficient
Avg. Area Income	21.540249
Avg. Area House Age	168198.049191
Avg. Area Number of Rooms	121591.585049
Avg. Area Number of Bedrooms	1573.129357
Area Population	15.229152

```
In [13]: prediction=model.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[13]: <matplotlib.collections.PathCollection at 0x1f2de399070>



```
In [14]: model.score(x_test,y_test)
```

Out[14]: 0.9169323598277197

```
In [56]: df2=pd.read_csv("9_bottle.csv")
df2
```

```
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3165:
DtypeWarning: Columns (47,73) have mixed types.Specify dtype option on import or s
et low_memory=False.
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

Out[56]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...
0	1	1	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0000A-3	0	10.500	33.4400	NaN	25.64900	NaN	...
1	1	2	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0008A-3	8	10.460	33.4400	NaN	25.65600	NaN	...
2	1	3	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0010A-7	10	10.460	33.4370	NaN	25.65400	NaN	...
3	1	4	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0019A-3	19	10.450	33.4200	NaN	25.64300	NaN	...
4	1	5	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0020A-7	20	10.450	33.4210	NaN	25.64300	NaN	...
...
864858	34404	864859	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0000A-7	0	18.744	33.4083	5.805	23.87055	108.74	...
864859	34404	864860	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0002A-3	2	18.744	33.4083	5.805	23.87072	108.74	...
864860	34404	864861	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0005A-3	5	18.692	33.4150	5.796	23.88911	108.46	...
864861	34404	864862	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0010A-3	10	18.161	33.4062	5.816	24.01426	107.74	...

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...
				20- 1611SR- MX-310- 2239- 09340264- 0015A-3							
	864862	34404	864863	093.4 026.4	15	17.533	33.3880	5.774	24.15297	105.66	...

864863 rows × 74 columns

In [57]: `df2.isna().sum()`

```
Out[57]: Cst_Cnt          0
Btl_Cnt          0
Sta_ID           0
Depth_ID         0
Depthm           0
          ...
TA1              862779
TA2              864629
pH2              864853
pH1              864779
DIC Quality Comment 864808
Length: 74, dtype: int64
```

In [58]: `df2.columns`

```
Out[58]: Index(['Cst_Cnt', 'Btl_Cnt', 'Sta_ID', 'Depth_ID', 'Depthm', 'T_degC',
'Salnty', 'O2ml_L', 'STheta', 'O2Sat', 'Oxy_μmol/Kg', 'BtlNum',
'RecInd', 'T_prec', 'T_qual', 'S_prec', 'S_qual', 'P_qual', 'O_qual',
'SThtaq', 'O2Satq', 'ChlorA', 'Chlqua', 'Phaeop', 'Phaqua', 'PO4uM',
'PO4q', 'SiO3uM', 'SiO3qu', 'NO2uM', 'NO2q', 'NO3uM', 'NO3q', 'NH3uM',
'NH3q', 'C14As1', 'C14A1p', 'C14A1q', 'C14As2', 'C14A2p', 'C14A2q',
'DarkAs', 'DarkAp', 'DarkAq', 'MeanAs', 'MeanAp', 'MeanAq', 'IncTim',
'LightP', 'R_Depth', 'R_TEMP', 'R_POTEMP', 'R_SALINITY', 'R_SIGMA',
'R_SVA', 'R_DYNHT', 'R_O2', 'R_O2Sat', 'R_SIO3', 'R_PO4', 'R_NO3',
'R_NO2', 'R_NH4', 'R_CHLA', 'R_PHAEO', 'R_PRES', 'R_SAMP', 'DIC1',
'DIC2', 'TA1', 'TA2', 'pH2', 'pH1', 'DIC Quality Comment'],
dtype='object')
```

In [59]: `df3=df2.drop(['DIC2', 'TA1', 'TA2', 'pH2', 'pH1', 'DIC Quality Comment'],axis=1)`

In []:

In []:

```
In [60]: df4=df3.drop(['BtlNum', 'T_qual', 'S_qual', 'O_qual', 'SThtaq', 'NH3uM', 'C14As1',  
                    'C14A1p', 'C14As2', 'C14A2p', 'DarkAs', 'DarkAp', 'MeanAs', 'MeanAp',  
                    'IncTim', 'LightP', 'R_NH4', 'R_SAMP', 'DIC1'],axis=1)
```

```
In [61]: df4.isna().sum()/len(df4)*100
```

```
Out[61]: Cst_Cnt      0.000000
        Btl_Cnt      0.000000
        Sta_ID       0.000000
        Depth_ID     0.000000
        Depthm       0.000000
        T_degC       1.267600
        Salnty       5.475318
        O2ml_L       19.501586
        STheta       6.092179
        O2Sat        23.540029
        Oxy_μmol/Kg   23.540723
        RecInd       0.000000
        T_prec       1.267600
        S_prec       5.475318
        P_qual       22.096910
        O2Satq       74.817168
        ChlorA       73.952869
        Chlqua       26.096272
        Phaeop       73.952984
        Phaqua       26.095809
        P04uM        52.210119
        P04q         47.762131
        Si03uM       59.058140
        Si03qu       40.930991
        NO2uM        60.967691
        NO2q         38.779437
        NO3uM        60.987694
        NO3q         38.726365
        NH3q         6.540227
        C14A1q       1.879835
        C14A2q       1.877754
        DarkAq       2.823915
        MeanAq       2.824031
        R_Depth      0.000000
        R_TEMP       1.267600
        R_POTEMP     5.324196
        R_SALINITY   5.475318
        R_SIGMA      6.111488
        R_SVA        6.101660
        R_DYNHT      5.394727
        R_O2         19.501586
        R_O2Sat      22.941784
        R_SI03       59.057215
        R_P04        52.209194
        R_NO3        60.986769
        R_NO2        60.966766
        R_CHLA       73.952406
        R_PHAEO      73.952522
        R_PRES       0.000000
dtype: float64
```

In []:

In [62]: df4.describe()

Out[62]:

	Cst_Cnt	Btl_Cnt	Depthm	T_degC	Salnty	O2ml_L	
count	864863.000000	864863.000000	864863.000000	853900.000000	817509.000000	696201.000000	812
mean	17138.790958	432432.000000	226.831951	10.799677	33.840350	3.392468	
std	10240.949817	249664.587267	316.050259	4.243825	0.461843	2.073256	
min	1.000000	1.000000	0.000000	1.440000	28.431000	-0.010000	
25%	8269.000000	216216.500000	46.000000	7.680000	33.488000	1.360000	
50%	16848.000000	432432.000000	125.000000	10.060000	33.863000	3.440000	
75%	26557.000000	648647.500000	300.000000	13.880000	34.196900	5.500000	
max	34404.000000	864863.000000	5351.000000	31.140000	37.034000	11.130000	

8 rows × 47 columns



In [63]: df4.shape

Out[63]: (864863, 49)

In [64]: df5=df4.iloc[0:5000,0:]
df5.shape

Out[64]: (5000, 49)

```
In [65]: per=df5.isna().sum()/len(df5)*100  
per1=pd.DataFrame(per,df5.columns,)  
per1
```

Out[65]:

	0
Cst_Cnt	0.00
Btl_Cnt	0.00
Sta_ID	0.00
Depth_ID	0.00
Depthm	0.00
T_degC	0.40
Salnty	3.04
O2ml_L	43.80
STheta	3.34
O2Sat	45.74
Oxy_μmol/Kg	45.74
Reclnd	0.00
T_prec	0.40
S_prec	3.04
P_qual	0.00
O2Satq	52.88
ChlorA	100.00
Chlqua	0.00
Phaeop	100.00
Phaqua	0.00
PO4uM	79.08
PO4q	20.92
SiO3uM	100.00
SiO3qu	0.00
NO2uM	100.00
NO2q	0.00
NO3uM	100.00
NO3q	0.00
NH3q	0.00
C14A1q	0.00
C14A2q	0.00
DarkAq	0.00
MeanAq	0.00
R_Depth	0.00
R_TEMP	0.40
R_POTEMP	4.50

	0
R_SALINITY	3.04
R_SIGMA	5.62
R_SVA	5.62
R_DYNHT	4.28
R_O2	43.80
R_O2Sat	46.20
R_SIO3	100.00
R_PO4	79.08
R_NO3	100.00
R_NO2	100.00
R_CHLA	100.00
R_PHAE0	100.00
R_PRES	0.00

```
In [66]: pr=per1[per1[0]>75].index  
pr
```

```
Out[66]: Index(['ChlorA', 'Phaeop', 'P04uM', 'Si03uM', 'NO2uM', 'NO3uM', 'R_SIO3',  
               'R_PO4', 'R_NO3', 'R_NO2', 'R_CHLA', 'R_PHAE0'],  
              dtype='object')
```

```
In [90]: df6=df5.drop(['ChlorA', 'Phaeop', 'PO4uM', 'SiO3uM', 'NO2uM', 'NO3uM', 'R_SI03',  
                    'R_PO4', 'R_NO3', 'R_NO2', 'R_CHLA', 'R_PHAE0'],axis=1)  
df6.isna().sum()/len(df5)*100
```

```
Out[90]: Cst_Cnt          0.00  
Btl_Cnt          0.00  
Sta_ID           0.00  
Depth_ID         0.00  
Depthm           0.00  
T_degC           0.40  
Salnty           3.04  
O2ml_L           43.80  
STheta           3.34  
O2Sat            45.74  
Oxy_μmol/Kg       45.74  
RecInd           0.00  
T_prec           0.40  
S_prec           3.04  
P_qual           0.00  
O2Satq           52.88  
Chlqua           0.00  
Phaqua           0.00  
P04q             20.92  
Si03qu           0.00  
NO2q             0.00  
NO3q             0.00  
NH3q             0.00  
C14A1q           0.00  
C14A2q           0.00  
DarkAq           0.00  
MeanAq           0.00  
R_Depth          0.00  
R_TEMP           0.40  
R_POTEMP         4.50  
R_SALINITY       3.04  
R_SIGMA          5.62  
R_SVA            5.62  
R_DYNHT          4.28  
R_O2             43.80  
R_O2Sat          46.20  
R_PRES           0.00  
dtype: float64
```


In [91]:

```
nul_col = df6.columns[df6.isnull().any()].tolist()
nul_col
```

Out[91]:

```
['T_degC',
 'Salnty',
 'O2ml_L',
 'STheta',
 'O2Sat',
 'Oxy_μmol/Kg',
 'T_prec',
 'S_prec',
 'O2Satq',
 'P04q',
 'R_TEMP',
 'R_POTEMP',
 'R_SALINITY',
 'R_SIGMA',
 'R_SVA',
 'R_DYNHT',
 'R_O2',
 'R_O2Sat']
```

In [88]:

```
#for i in nul_col:
    #df6[i]= df6.fillna(df6[i].mean())
```

In [89]: *#df6*

Out[89]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_De
0	1	1	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0000A-3	0	1	1	1	1	1	...	
1	1	2	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0008A-3	8	1	1	1	1	1	...	
2	1	3	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0010A-7	10	1	1	1	1	1	...	1
3	1	4	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0019A-3	19	1	1	1	1	1	...	1
4	1	5	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0020A-7	20	1	1	1	1	1	...	2
...	
4995	165	4996	092.0 098.0	19- 4904NS- HY-102- 1342- 09200980- 0099A-3	99	165	165	165	165	165	...	9
4996	165	4997	092.0 098.0	19- 4904NS- HY-102- 1342- 09200980- 0100A-7	100	165	165	165	165	165	...	10
4997	165	4998	092.0 098.0	19- 4904NS- HY-102- 1342- 09200980- 0125A-7	125	165	165	165	165	165	...	12
4998	165	4999	092.0 098.0	19- 4904NS- HY-102- 1342- 09200980- 0149A-3	149	165	165	165	165	165	...	14

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_De
4999	165	5000	19-4904NS-HY-102-1342-09200980-0150A-7	150	165	165	165	165	165	...	15

5000 rows × 37 columns


```
In [94]: df6.fillna(df6["T_degC"].mean())
df6.fillna(df6["Salnty"].mean())
df6.fillna(df6["O2ml_L"].mean())
df6.fillna(df6["STheta"].mean())
df6.fillna(df6["O2Sat"].median())
df6.fillna(df6["Oxy_μmol/Kg"].mean())
df6.fillna(df6["T_prec"].mean())
df6.fillna(df6["S_prec"].mean())
df6.fillna(df6["O2Satq"].mean())
df6.fillna(df6["PO4q"].mean())
df6.fillna(df6["R_TEMP"].mean())
df6.fillna(df6["R_POTEMP"].mean())
df6.fillna(df6["R_SALINITY"].mean())
df6.fillna(df6["R_SIGMA"].mean())
df6.fillna(df6["R_SVA"].mean())
df6.fillna(df6["R_DYNHT"].mean())
df6.fillna(df6["R_O2"].mean())
df6=df6.fillna(df6["R_O2Sat"].mean())
```

```
In [95]: df6.isna().sum()
```

```
Out[95]: Cst_Cnt      0
         Btl_Cnt      0
         Sta_ID      0
         Depth_ID     0
         Depthm      0
         T_degC      0
         Salnty      0
         O2ml_L      0
         STheta      0
         O2Sat      0
         Oxy_μmol/Kg  0
         RecInd      0
         T_prec      0
         S_prec      0
         P_qual      0
         O2Satq      0
         Chlqua      0
         Phaqua      0
         PO4q        0
         SiO3qu      0
         NO2q        0
         NO3q        0
         NH3q        0
         C14A1q      0
         C14A2q      0
         DarkAq      0
         MeanAq      0
         R_Depth     0
         R_TEMP      0
         R_POTEMP    0
         R_SALINITY  0
         R_SIGMA     0
         R_SVA       0
         R_DYNHT     0
         R_O2        0
         R_O2Sat     0
         R_PREX      0
         dtype: int64
```


```
In [98]: df7=df6.iloc[:10,:10]  
sns.pairplot(df7)
```

```
Out[98]: <seaborn.axisgrid.PairGrid at 0x1f287b2eac0>
```



```
In [97]: sns.heatmap(df7.corr())
```

```
Out[97]: <AxesSubplot:>
```



```
In [107]: df6=df6.drop(["Sta_ID", "Depth_ID"],axis=1)
```

```
In [109]: x=df6.drop(["R_O2"],axis=1)
y=df6["R_O2"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
model1=LinearRegression()
model1.fit(x_train,y_train)
model1.intercept_
```

```
Out[109]: -4.618527782440651e-14
```


```
In [110]: coeff=pd.DataFrame(model1.coef_,x.columns,columns=["Coefficient"])  
coeff
```


Out[110]:

	Coefficient
Cst_Cnt	-1.900857e-14
Btl_Cnt	5.568462e-16
Depthm	2.371912e-15
T_degC	7.851962e-16
Salnty	1.374703e-14
O2ml_L	1.000000e+00
STheta	-9.098279e-15
O2Sat	2.150917e-15
Oxy_μmol/Kg	7.801648e-16
Reclnd	3.879562e-16
T_prec	3.818522e-15
S_prec	-3.387600e-15
P_qual	-1.942890e-16
O2Satq	-3.325655e-16
Chlqua	2.775558e-17
Phaqua	1.387779e-17
PO4q	1.632673e-16
SiO3qu	0.000000e+00
NO2q	-1.734723e-18
NO3q	0.000000e+00
NH3q	0.000000e+00
C14A1q	0.000000e+00
C14A2q	0.000000e+00
DarkAq	0.000000e+00
MeanAq	0.000000e+00
R_Depth	2.433526e-15
R_TEMP	1.450847e-16
R_POTEMP	9.528765e-16
R_SALINITY	1.280675e-14
R_SIGMA	-1.068890e-15
R_SVA	-2.690356e-15
R_DYNHT	1.852546e-16
R_O2Sat	1.030526e-16
R_PRES	-5.059402e-15

```
In [113]: pred=model1.predict(x_test)  
plt.scatter(y_test,pred)
```

Out[113]: <matplotlib.collections.PathCollection at 0x1f28a3286d0>



```
In [116]: model1.score(x_test,y_test)
```

Out[116]: 1.0